



CS 5594: BLOCKCHAIN TECHNOLOGIES

Spring 2021

THANG HOANG, PhD

DISTRIBUTED SYSTEMS

Outline

Definition and Characteristics

Architecture Models

Distributed Algorithms

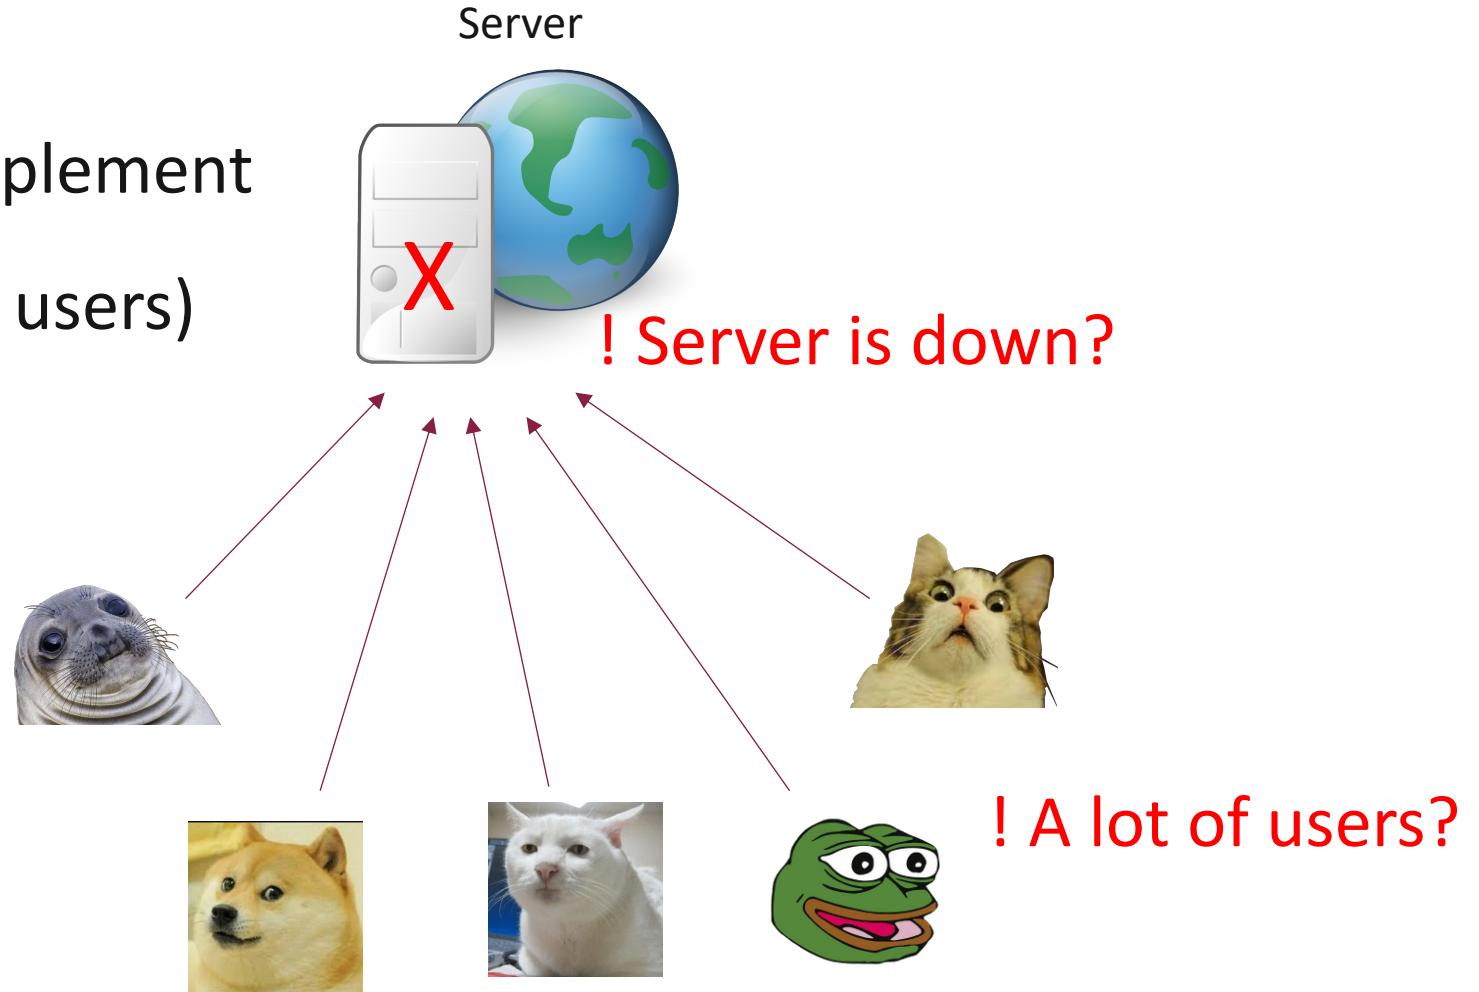
Distributed Systems

Definition and Characteristics

Centralized Systems

All tasks completed by a **single** entity

- ✓ Simple
- ✓ Easy to design and implement
- ✓ Efficient (with small # users)



Distributed Systems

A group of independent entities communicated with one another in a coordinated manner

Collaboratively enable a service (computing, data sharing and storage)

Goal: Address inherent limitations of centralized systems

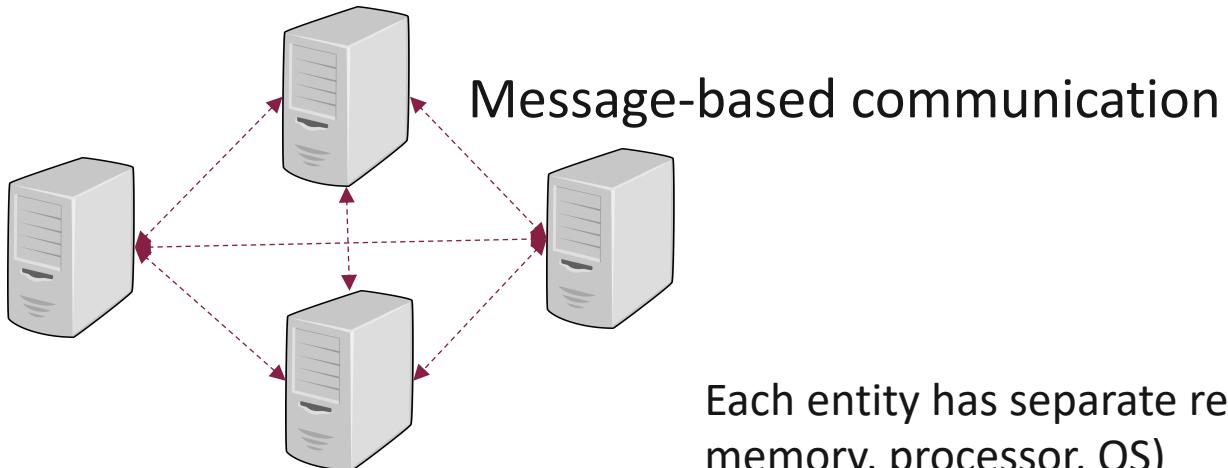
Robustness

Scalability

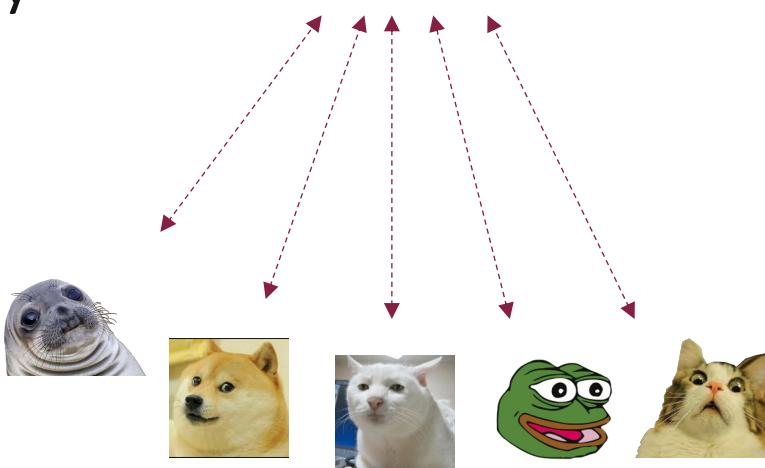
Reliability

Distributed Systems

Distributed server



Appeared as a **single** entity



Key Characteristics

Transparency

Most important feature

Illusion of a single system

Hide all internal organization, communication details

Uniform interface

Access transparency, location transparency, relocation transparency, migration transparency, replication transparency, concurrency transparency, failure transparency, scaling transparency, performance transparency

Key Characteristics

Openness

Heterogeneity

Variety and differences in hardware and software components

Resource Sharing

Resources (hardware, software, data) accessed across multiple entities

Concurrency

Parallel executions of activities

Reduce latency, increase throughput

Key Characteristics

Scalability

Add/remove components to/from the system

Fault Tolerance

Continuous availability

Design Goals

High Performance

Low latency, high throughput

Reliability

Preserve **correctness** and **integrity** in the presence of faulty/malicious nodes

Failure detection, self-stabilization

Scalability

Adapt with flexible number of users in the system

Design Goals

Consistency

Update consistency, replication consistency, cache consistency, failure consistency, clock consistency, user interface consistency

Synchronization between concurrent tasks

Security

Malicious adversaries, secure communication, resource protection

Impossibility Result

CAP Theorem

“Any distributed system cannot achieve Consistency, Availability and Partition tolerance concurrently.”

Gilbert and Lynch

Impossibility Result

Consistency

All nodes see the same data at the same time

Availability

If the node in the system does not fail, it must always respond to the user's request.

Partition tolerance

The network will be allowed to lose arbitrarily many messages sent from one node to another

Choose 2 out of 3

Generally between consistency & availability under partition

Distributed System Applications

Distributed systems are everywhere

Mobile systems

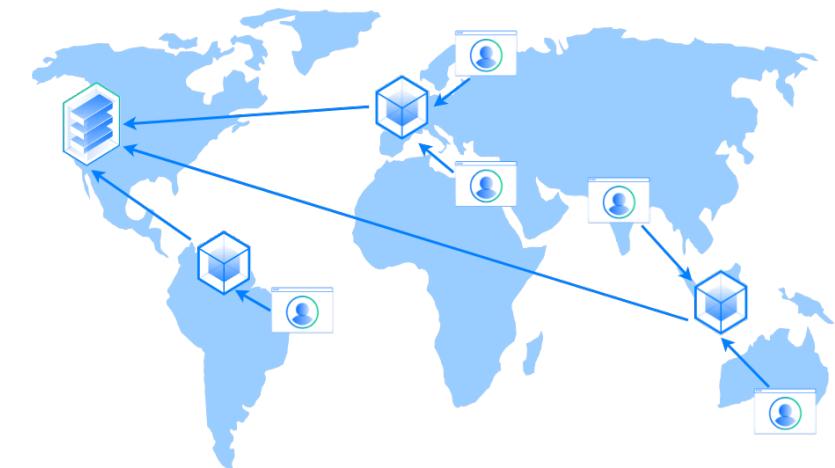
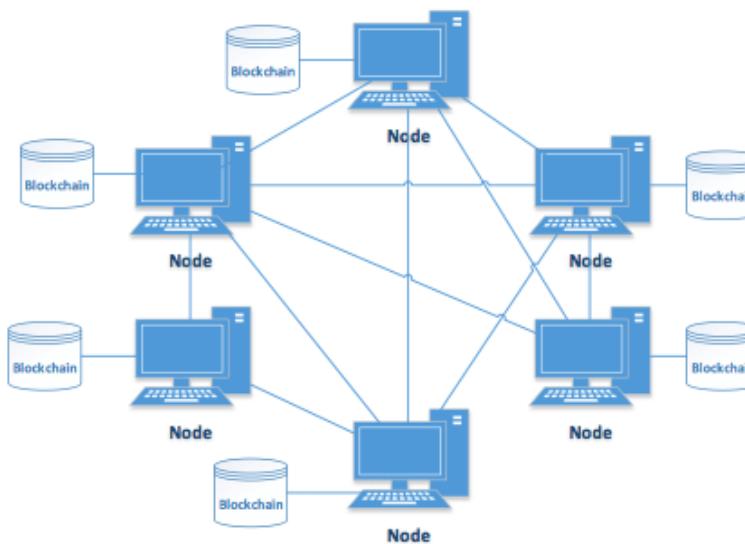
Sensor networks

IoT

Ubiquitous and Pervasive computing

WWW

P2P computing



Content Delivery Network (CDN)

Distributed Systems

Architectural Models

Client-Server Architecture

Basic model

Two types of node: client (slave) and server (master)

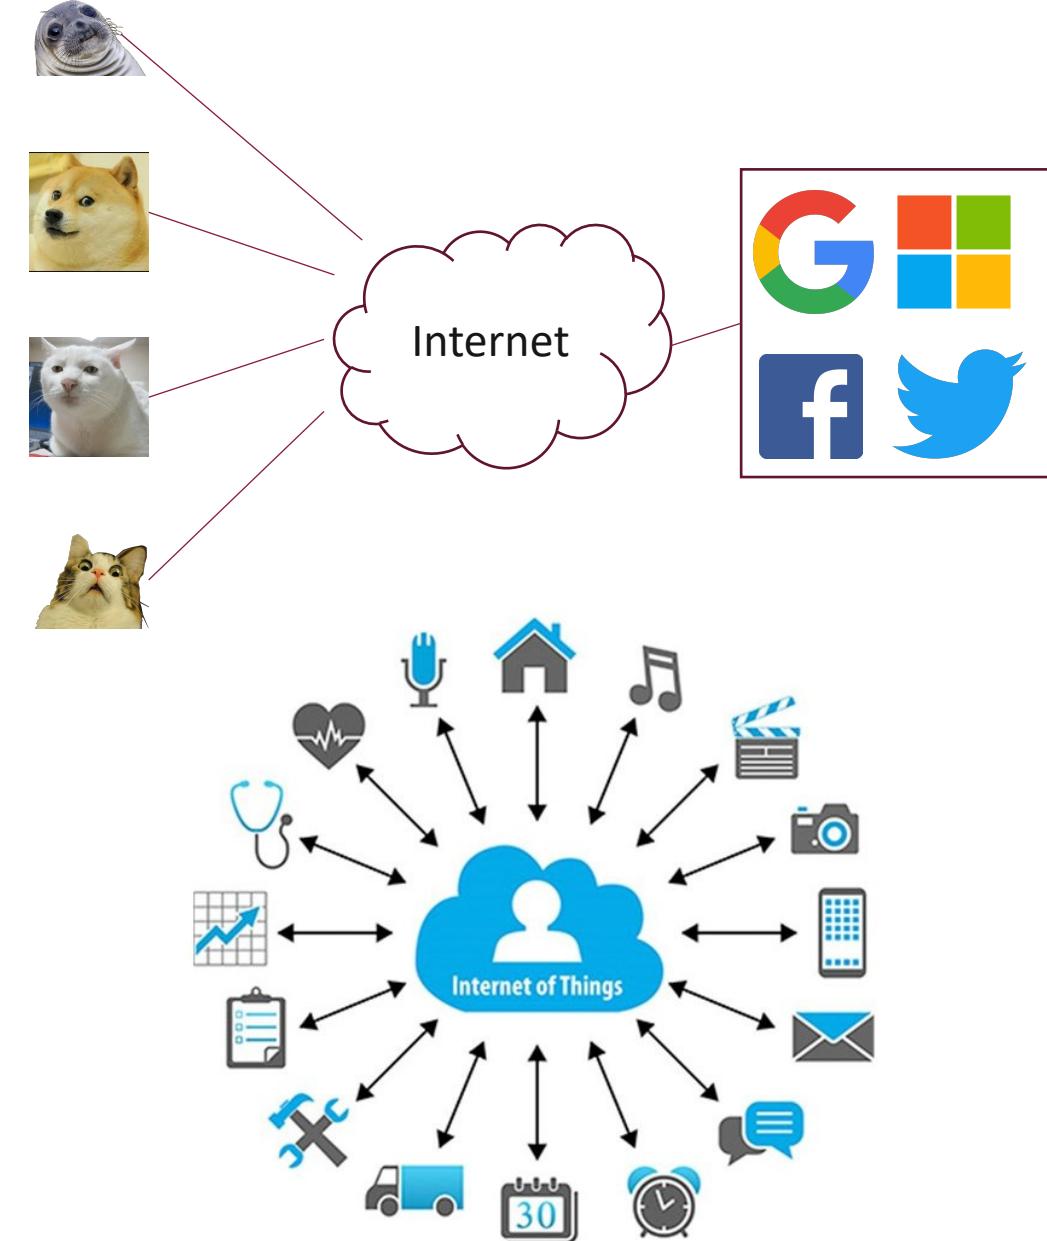
All tasks accomplished by server

Server is resource-powerful

Client is resource-limited

Asymmetric, partially distributed

Examples: Cloud services (Amazon, MS, Facebook, Google), IoTs



Client-Server Architecture

Advantage

Easy to maintain security and reliability

Enable a wide range of services

Easy to design and implement

Client-Server Architecture

Disadvantage

Central point of failure and compromise

Attacks targeting to server nodes (e.g., DoS, data-breach)

Resource management and administration

Central point of trust

Server has more control and authority in the system

Not so scalable

More clients join, more server demands

Peer-to-Peer Architecture

A network of nodes (peers) sharing resources directly with each other

Symmetric: All nodes are equal participants and play both roles: provider and consumer of resource

No *server* node

Fully distributed, no centralized data and resource

“The ultimate form of democracy on the Internet”

Examples: blockchains, vehicular network,
file-sharing

Bitcoin: A Peer-to-Peer Electronic Cash System

Satoshi Nakamoto
satoshin@gmx.com
www.bitcoin.org

Abstract. A purely peer-to-peer version of electronic cash would allow online payments to be sent directly from one party to another without going through a financial institution. Digital signatures provide part of the solution, but the main benefits are lost if a trusted third party is still required to prevent double-spending. We propose a solution to the double-spending problem using a peer-to-peer network.

Peer-to-Peer Architecture

Advantage

Distributed trust

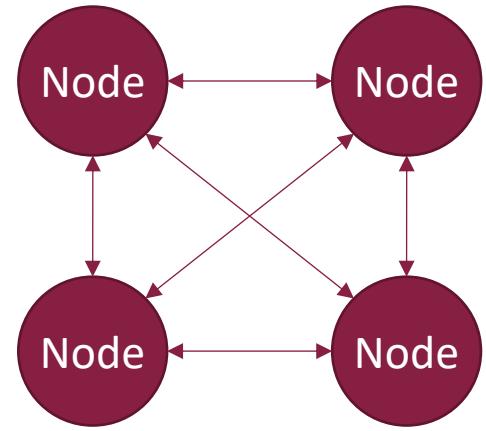
Balanced resource load

High resource capacity and high scalability

More clients, more servers

High fault-tolerance and resiliency against

DoS attacks



Peer-to-Peer Architecture

Disadvantage

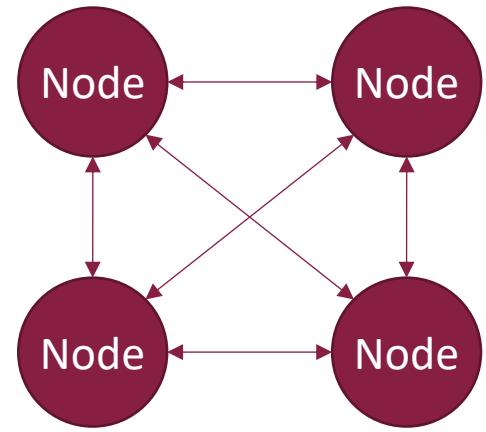
Costly backup, high bandwidth consumption

Hard to control

Hard to maintain security and consistency

Vulnerable to network partitions,
byzantine behavior

Unstable



Distributed vs. Decentralized

P2P is distributed, but offers various degrees of decentralization

Some P2P still need central authorities to make decision (e.g., network control, resource load) efficiently

Somewhat centralized

Decentralized is NOT all-or-nothing

Distributed vs. Decentralized

In fact, no system is purely decentralized, or purely centralized

Blockchain can be centralized or decentralized under certain degrees

Depend on the design and application requirements

Unstructured P2P network

Easy to build

Loose restriction on overlay structure, data location and resource distribution

Nodes communicate randomly, perform arbitrary tasks

High resiliency to **churn**

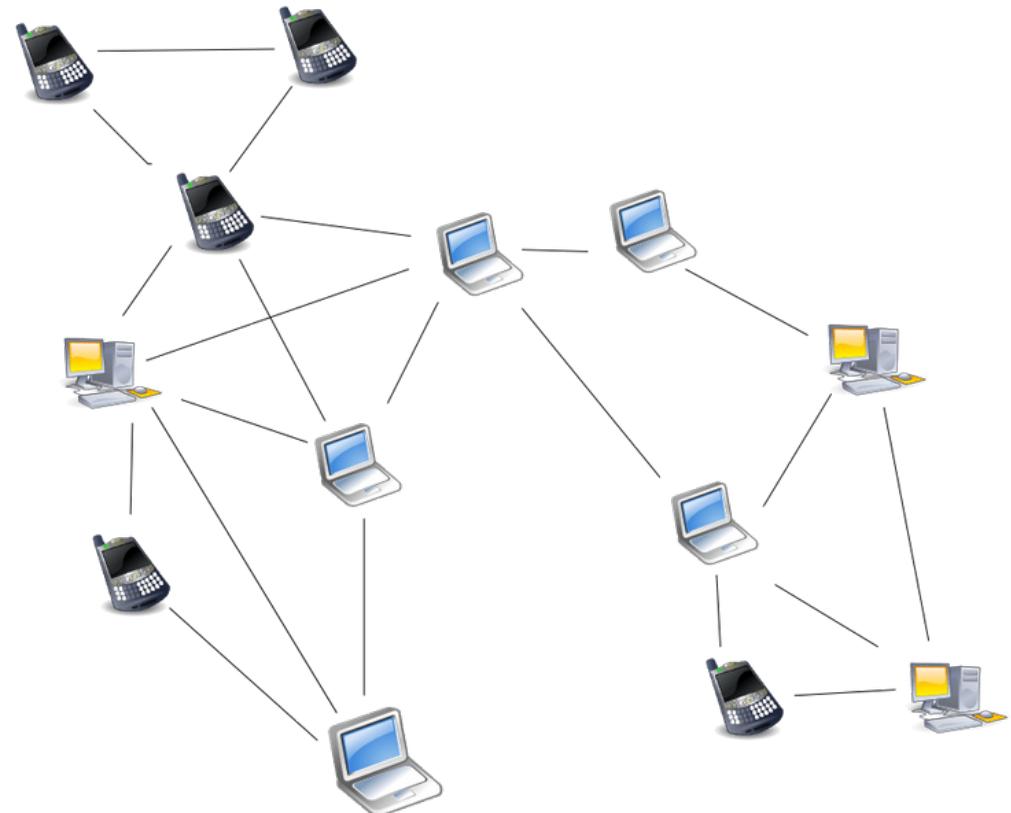
Nodes leave and join frequently

Nodes and resources are loosely-coupled

Data navigation issue

High resource (CPU, memory, network) usage

Example: Napster, Gnutella, KaZaA



Structured P2P network

Structured overlay network, restriction on content placement and resource distribution

Nodes and resources are **tightly-coupled**, everyone has their own task

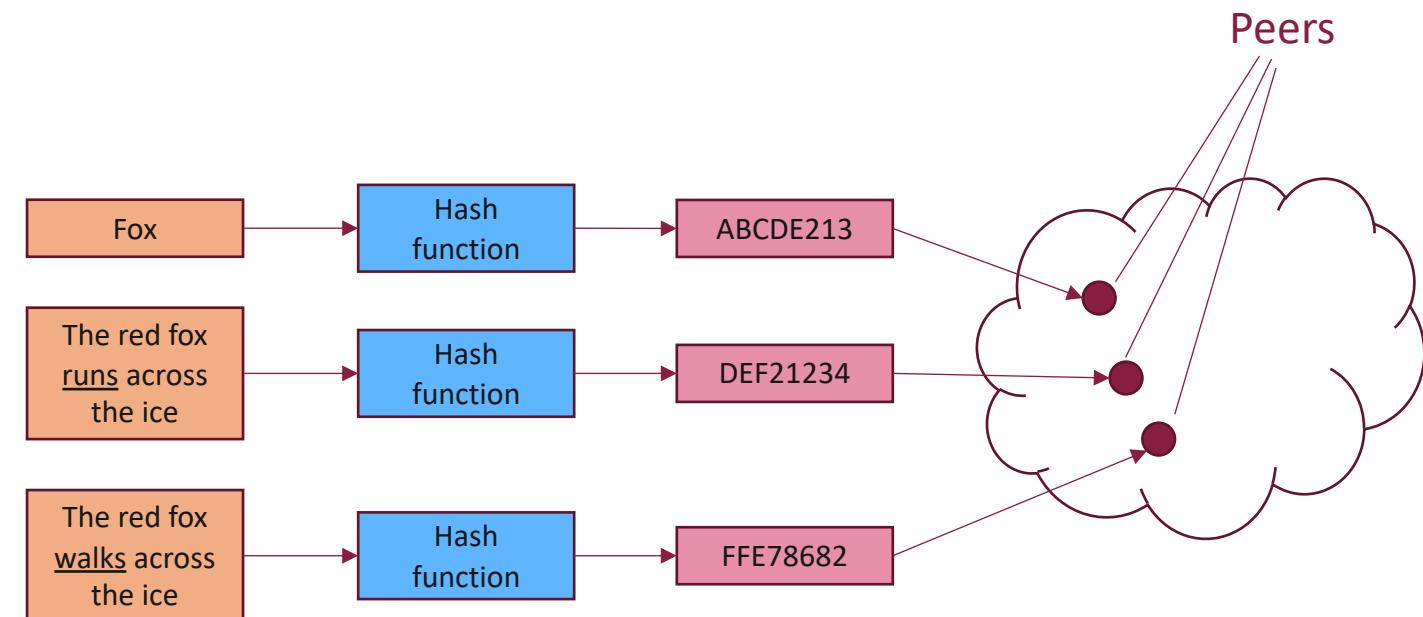
Each node is responsible for a specific role in the network

Distributed Hash Table (DHT) for node-task assignment

Simplifying content location

Harder to build

Low resiliency against churn



Hybrid P2P network

Central authorities to help nodes navigate each other

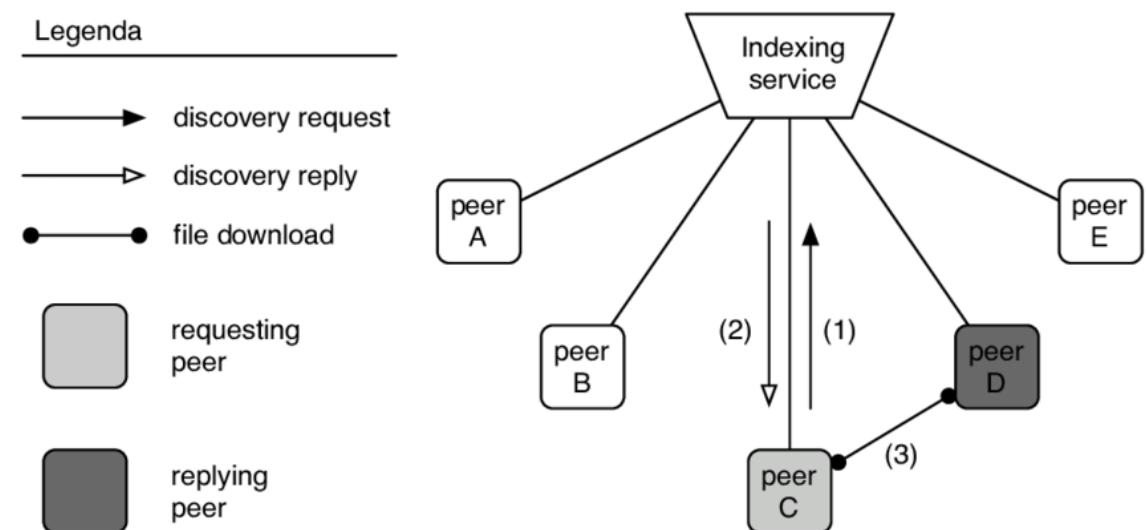
Combine client-server with P2P models

Tend to improve overall performance

Trade-off b/w centralization vs. node equality

Inherit the best of both worlds

Efficiency in C-S setting, and decentralization in P2P setting



Distributed Algorithms

Consensus Mechanism

Consensus Mechanism

Main Motivation: Reliability and Fault-Tolerance in distributed system

Correct operation in the presence of corrupted nodes

Reach a common agreement in a distributed/
decentralized system

Nodes propose values

All nodes must agree on one of these values

Consensus Mechanism

Key to solving many problems in distributed computing

Atomic broadcast

Atomic commit of database transaction

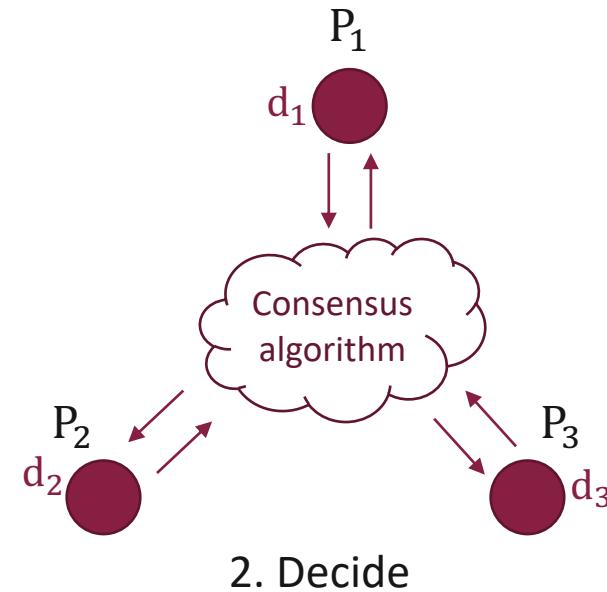
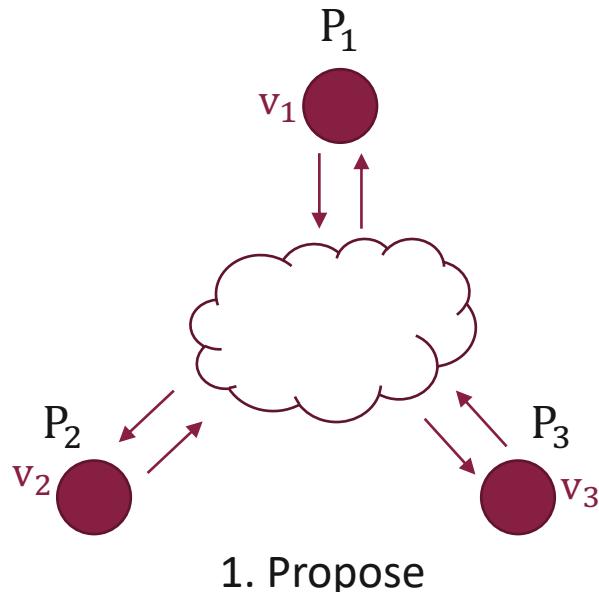
Clock synchronization

Dynamic group membership

Consensus Protocol: Definition

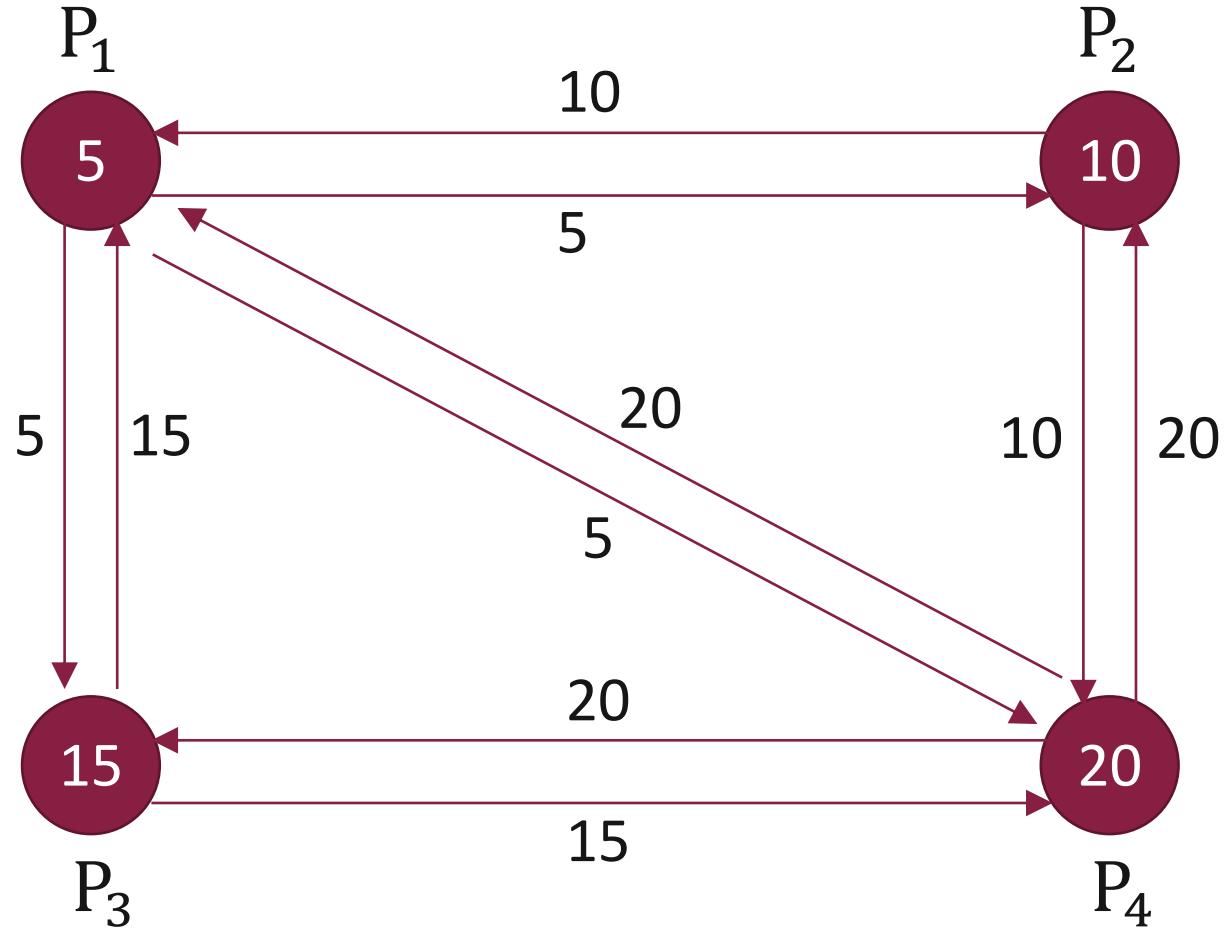
A consensus protocol comprises two algorithms:

- $v_i \leftarrow \text{Propose}()$: Each node n_i propose a value v_i and broadcast v_i to the network
- $v \leftarrow \text{Decide}(v_1, \dots, v_n)$: All nodes agree a common value $v \in \{v_1, \dots, v_n\}$
- The protocol **terminates** when all correct nodes decide on the **same value**
- The agreed value **cannot** be arbitrary: it must come from some correct node



Consensus Protocol

Example: Find max value among all values



Consensus Properties

Validity

Value agreed is a value **proposed**

Agreement

All **correct** nodes agree on the **same** value

Integrity

Every **correct** node decides at most **once**

Consensus Properties

Termination

Every **correct** node must decide at the end of protocol

Safety

Every **correct** node must not agree on **incorrect** value

Liveness

Every **correct** value must be **accepted**

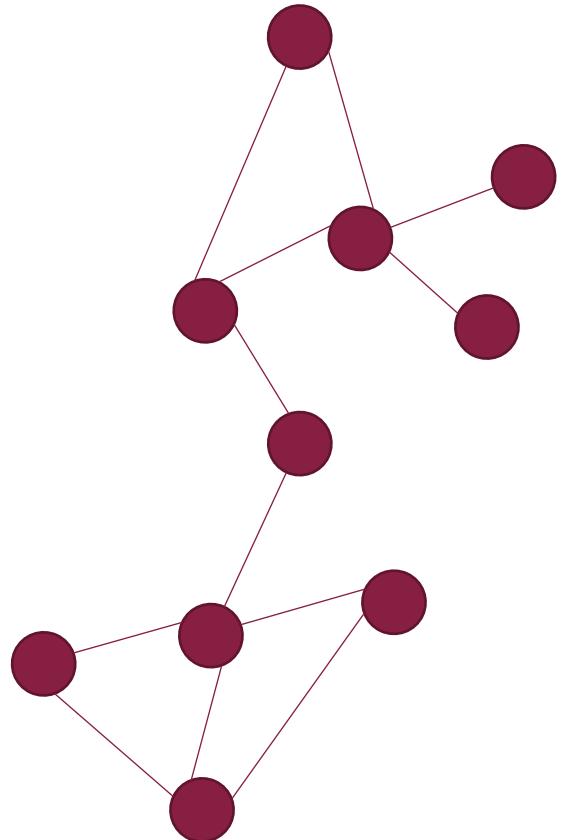
When Failure Happens

If no failure or malice, easy to reach a consensus

Individuals broadcast their values to all nodes

Values received with a pre-defined timeframe (synchronous)

What if there are failures or malicious activities in the network?



When Failure Happens

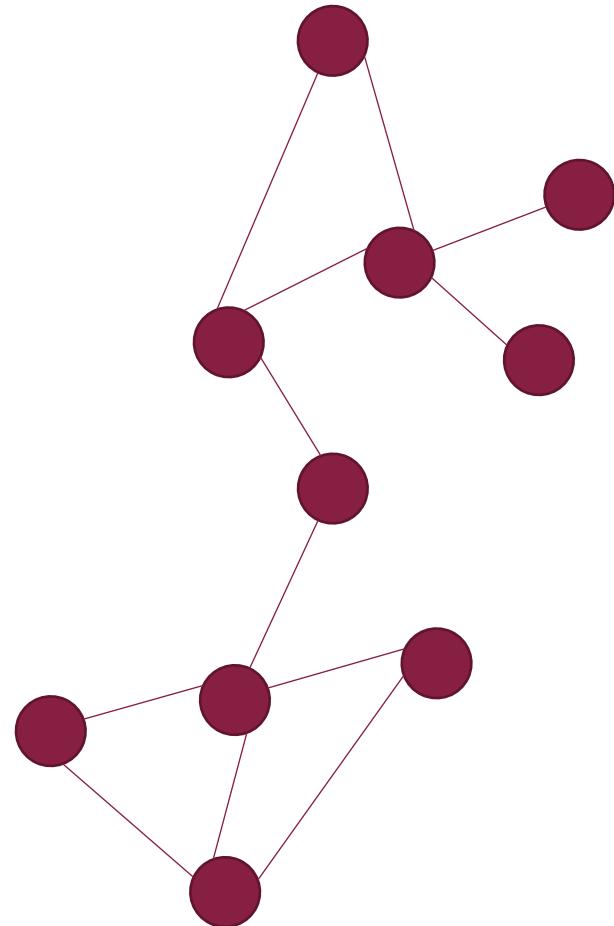
Common types of failure

Crash Fault: Node crashed, offline during communication

Network Fault: Not all pairs of nodes well-connected (partitioned network), latency (no notion of global time)

Byzantine Fault: Nodes may be malicious

Achieving consensus in the faulty (yet realistic) environment is **hard**



Synchronous vs. Asynchronous Systems

Synchronous system

Defined maximum waiting time for message transmission

Easy to reach a consensus

Asynchronous system

Undefined waiting time

Hard to achieve a consensus

Impossibility Results

(Another) Impossibility Results

Byzantine Generals Problem

“Consensus is impossible with a single faulty node”

- Fischer, Lynch and Patterson

Choose 2 out of 3: Safety, Liveness, Fault-Tolerance

Impossibility Results

Understand IRs correctly

IRs are more about the **model** than about the problem

Developed to study systems like distributed databases

Blockchain has different model

Consensus is still useful and achievable

Find right algorithm for specific application domain

Consensus Algorithms

Paxos

Majority rule, asynchronous setting

Consistency, fault-tolerance, but may get stuck (2 out of 3 rule)

Byzantine-fault **intolerance**

Raft

Leader-Follower model

Choose 2 in 3: Safety, Liveness, Fault-Tolerance

Byzantine-fault **intolerance**

<http://thesecretlivesofdata.com/raft/> (animated example)

BFT

Byzantine-fault tolerance

Stay tuned for next lecturers!

Consensus in Public Blockchain

Traditional consensus works on **closed** environment

Nodes know addresses of their peers

Every node accesses a shared memory

Public blockchain is an **open** P2P system

Where to keep shared memory in P2P?

Anyone can join and leave the network at anytime

How to enable consensus in an open system?



Signed by Bob
Pay to pk_{Bob} : $H()$



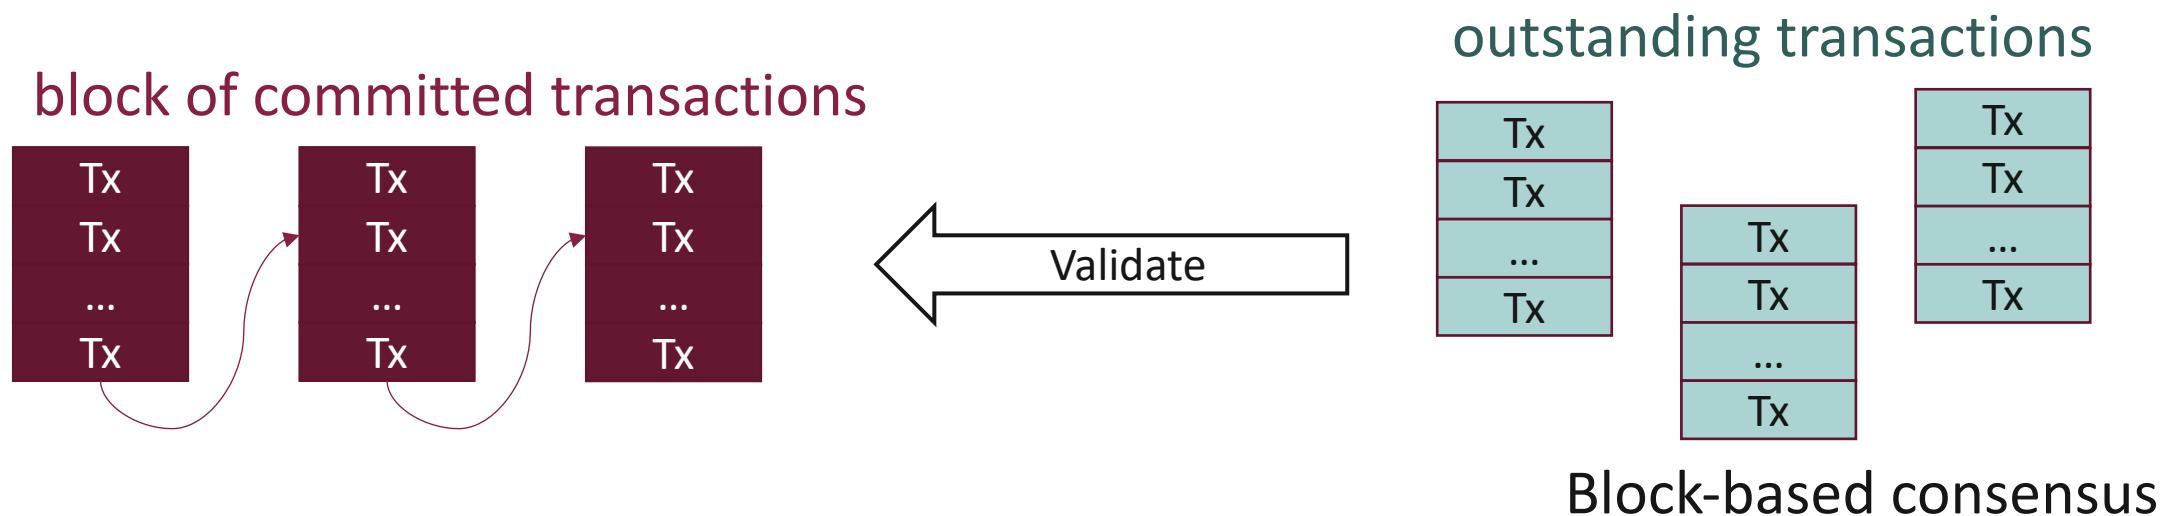
All nodes must agree on the validity of the Bob's transaction

Consensus in Public Blockchain

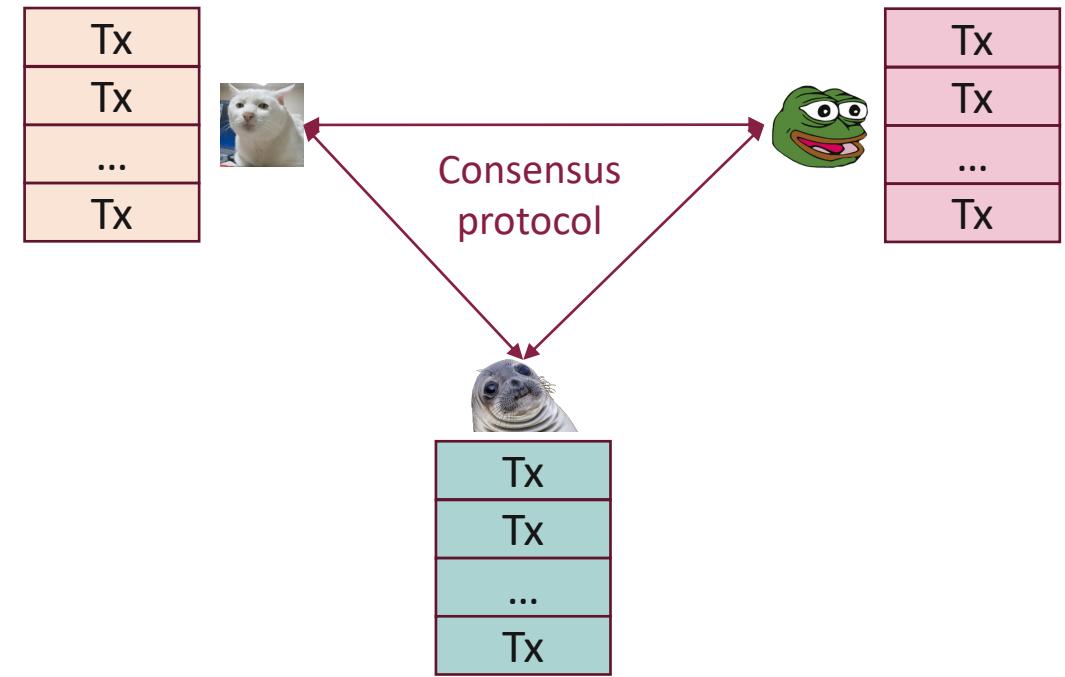
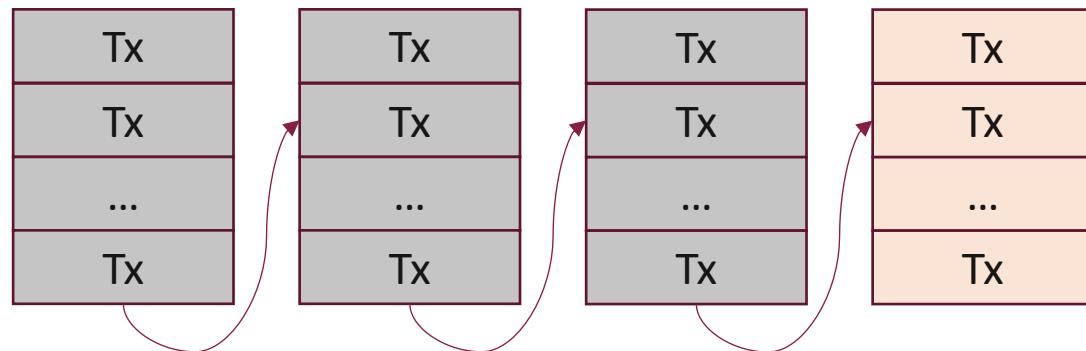
At any given time:

All nodes have a sequence of blocks of transactions
they have reached a consensus on (**block of committed transactions**)

Each node has a set of **outstanding transactions**
that need to be validated against **block of committed transactions**



Blockchain Consensus



Consensus in Public Blockchain

Bitcoin introduces incentive concept for honest actions

Possible as Bitcoin is a digital currency

Embrace randomness

Does away with the notion of a specific end-point

Consensus happens over long-time scales – approx. 1 hour

Blockchain consensus works better in practice than in theory

Theory is catching up

Theory is still very important as It can help predict unforeseen attacks

We will find out in more details later...